


# **MAXGBoost: A Fast Novel Heuristic Approach to Adaptive Learning Rates in Gradient Boosted Decision Trees**

By: Anirudh Chintaluri, Radin Rezanezhad


# Fundamentals

- GBDTs initialize with a decision tree that's evaluated using a loss function
- The model minimizes loss by moving in the direction of the negative gradient
- This process is similar to neural networks' backpropagation, but creates an ensemble decision tree model

A Gradient Boosting Decision Trees Structure

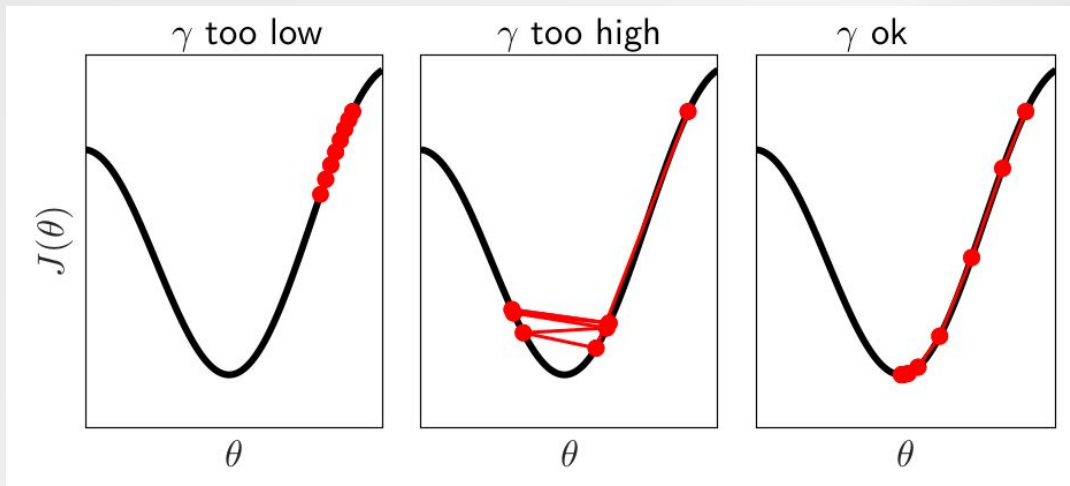


A Decision Trees Structure





# Learning Rate Mechanics

- The algorithm adds new decision trees scaled by a learning rate ( $\eta$ )
- Learning rate determines how far down the loss function the tree will go
- Low  $\eta$  ensures convergence but requires more iterations and computational cost
- High  $\eta$  speeds up convergence but risks overshooting the optimal solution




# Learning Rate Optimization





- Optimal learning rate depends on both dataset and loss function
- Tuning requires iterative experimentation
- Adaptive strategies help reduce need for extensive hyperparameter searches
- Particularly important for datasets with high sparsity



# Introduction to MAXGBoost

- Momentum Approximation  
XGBoost dynamically adjusts learning rate based on loss momentum
- Enables faster convergence
- Reduces need for manual optimization while maintaining high detection accuracy



- |   |                                      |
|---|--------------------------------------|
| A   | Starting point                       |
| B   | Optimal point of loss function       |
|  | GD                                   |
|  | MGD                                  |
|  | Gradient component of each iteration |
|  | Momentum component of each iteration |

# The Dataset

- Credit Card Fraud Detection
- 284,807 total instances
- 31 total attributes
- Notable Features:
  - Time
  - Amount
  - V1-V38 → Result of PCA
- 492 instances marked as fraud
  - < 0.2% of data
- Dataset already preprocessed



# The Algorithm - Traditional Neural Network

$$w_{t+1} = w_t - \eta \nabla L$$

# The Algorithm - Neural Network with Momentum

- $v$  = velocity
- $\beta$  = how influential past velocities are on the new velocity calculated
- Converges faster than with a constant learning rate

$$v_{t+1} = \beta v_t + (1 - \beta) \nabla L$$

$$w_{t+1} = w_t - \eta v_{t+1}$$



# The Algorithm - MAXGBoost

- Combines GBDT and momentum-based updates
- Approximates  $\nabla L$
- Using momentum to update the **learning rate**
  - Loss increase =  $\eta$  increase
  - Loss decrease =  $\eta$  decrease

$$v_{t+1} = \beta v_t + (1 - \beta)(L_{t+1} - L_t)$$

$$\eta_{t+1} = \eta_t(1 + v_{t+1})$$

# Experiments - Dataset

## Splits

- 68%: Training
- 12%: Validation
- 20%: Testing

## Analysis

- 5-fold Cross-validation
- Results based on best fold per model

# Experiments - Models

<b>Decision Tree</b>	<b>Single Decision Tree</b>
<b>Random Forest</b>	<b>Ensemble Voting Decision Tree Model</b>
<b>Constant <math>\eta</math> XGBoost</b>	<b>XGBoost with a constant learning rate</b>
<b>Exponential Decay XGBoost</b>	<b>XGBoost with a learning rate multiplied by constant every iteration</b>
<b>MAXGBoost</b>	<b>XGBoost with Momentum Approximation</b>

# Experiments - Hyperparameters

<b>Decision Tree</b>	<b>Max Depth = 4</b>
<b>Random Forest</b>	<b>Estimators = 5</b>
<b>Constant <math>\eta</math> XGBoost</b>	<b>Estimators = 423, <math>\eta</math> = 0.08</b>
<b>Exponential Decay XGBoost</b>	<b>Estimators = 423, <math>\alpha</math> = 0.9</b>
<b>MAXGBoost</b>	<b>Estimators = 423, Initial <math>\eta</math> = 0.89, <math>\beta</math> = 0.99</b>

# Results

Model	Accuracy	Precision	Recall	AUC
Decision Tree	0.99946	0.86813	0.80612	0.90296
Random Forest	0.99951	0.77000	0.93902	0.96931
Constant $\eta$	0.99977	0.88636	0.96296	0.98139
Exponential decay	0.99979	0.88764	<b>0.97531</b>	<b>0.98757</b>
MAXGBoost	<b>0.99980</b>	<b>0.93827</b>	0.92683	0.96337

# Discussion - MAXGBoost Strengths

- **Best Precision** in Fraud Detection
  - Exceptional performance in minimizing false positives
  - Effective for high-stakes financial environments
    - Reduces customer friction from false alerts
    - Minimizes operational costs from investigation overhead

## Discussion - MAXGBoost Limitations

- Relatively low AUC for Fraud Detection
  - **Rapid** learning rate **decrease** with small loss improvements
  - **Limited exploration of feature space**
  - Over-emphasis on strongly discriminative features
  - Underweighting of subtle fraud patterns

# Conclusion

- Decision Tree-based models effective in classifying positive class of fraud **in < 0.2% of data**
  - **Increases the ability for law enforcement** to correctly catch frauds
- MAXGBoost best model comparing accuracy and precision
  - MAXGBoost most effective in **preventing false accusations of fraud** and **increasing customer satisfaction**
- Exponential decay XGBoost most effective in **catching fraudulent cases**, compromising false accusations



# Future Work

- **Hybrid Model Development**
  - Combine strengths of EGB and MAXGBoost
  - Create integrated momentum-decay approach
- Feature-specific momentum updates
- Class-aware momentum adjustments

# Thanks!

Questions?

# References

- Andrej Baranovskij. (2019, March 12). *Selecting Optimal Parameters for XGBoost Model Training*. Medium; Towards Data Science. <https://medium.com/towards-data-science/selecting-optimal-parameters-for-xgboost-model-training-c7cd9ed5e45e>
- Beja-Battais, P. (2023, October 6). *Overview of AdaBoost: Reconciling its views to better understand its dynamics*. ArXiv.org. <https://doi.org/10.48550/arXiv.2310.18323>
- Brownlee, J. (2016, September 15). *Tune Learning Rate for Gradient Boosting with XGBoost in Python*. Machine Learning Mastery. <https://machinelearningmastery.com/tune-learning-rate-for-gradient-boosting-with-xgboost-in-python/>
- Chen, T., & Guestrin, C. (2016). XGBoost: a Scalable Tree Boosting System. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining - KDD '16*, 785–794. <https://doi.org/10.1145/2939672.2939785>
- Credit Card Fraud Detection. (n.d.). Kaggle. <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud/data>
- Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., & Liu, T.-Y. (2017). *LightGBM: A Highly Efficient Gradient Boosting Decision Tree*. Neural Information Processing Systems; Curran Associates, Inc. [https://papers.nips.cc/paper\\_files/paper/2017/hash/6449f44a102fde848669bdd9eb6b76fa-Abstract.html](https://papers.nips.cc/paper_files/paper/2017/hash/6449f44a102fde848669bdd9eb6b76fa-Abstract.html)
- Learning rate. (n.d.). Wikipedia. [https://en.wikipedia.org/wiki/Learning\\_rate](https://en.wikipedia.org/wiki/Learning_rate)
- Mohan, A. (2021, April 6). *IEEE-CIS Fraud Detection - Top 5% Solution - Towards Data Science*. Medium; Towards Data Science. <https://medium.com/towards-data-science/ieee-cis-fraud-detection-top-5-solution-5488fc66e95f>
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Müller, A., Nothman, J., Louppe, G., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., & Duchesnay, É. (2018). Scikit-learn: Machine Learning in Python. *ArXiv:1201.0490 [Cs]*. <https://arxiv.org/abs/1201.0490>
- Vitaly Bushaev. (2017, December 4). *Stochastic Gradient Descent with momentum - Towards Data Science*. Medium; Towards Data Science. <https://medium.com/towards-data-science/stochastic-gradient-descent-with-momentum-a84097641a5d>
- Wang, C., Wang, Z., Ouyang, Y., & Soleimani, B. H. (2024, May 27). *Adaptive Learning Rates for Gradient Boosting Machines*. Proceedings of the Canadian Conference on Artificial Intelligence. <https://caiac.pubpub.org/pub/py65wd3c/release/1>